

AI Agent Orchestration

Concept Design Document

Secure, Identity-Aware AI Agent Execution for the Enterprise

Seed Stage | Confidential | 2025

Executive Summary

 Aligns with: Slide 1 – Cover

AI Agent Orchestration is a security-first SaaS platform purpose-built to bring enterprise-grade control, governance, and identity enforcement to autonomous AI agent execution. As organizations race to deploy AI across operational workflows, security and compliance teams are blocking production rollouts due to a fundamental gap: existing tools offer agent capability without operational trust.

This document outlines the concept, architecture, market positioning, and go-to-market strategy for a platform that closes that gap — delivering a centralized control plane that orchestrates AI agents with the rigor Fortune 500 enterprises require.

Fortune 500 & Regulated Enterprises Target Market	SaaS Subscription + Usage Model	Seed Fundraise Stage
---	---	--------------------------------

The Market Problem

 Aligns with: Slide 2 – Market Problem

AI agent demand is outpacing enterprise trust and control frameworks. The friction is not lack of interest — enterprises are eager to deploy autonomous execution across operational workflows. The friction is confidence: security teams will not approve production deployment when control guarantees are weak or undefined.

Core Market Gaps

- Enterprises want autonomous execution across operational workflows
- Security teams block production deployment due to weak control guarantees

- Policy-only approaches fail under compliance and audit scrutiny
 - The market lacks an enterprise-grade execution platform — not just agent UX
-

Why Now

 Aligns with: Slide 3 – Urgency

The cost of inaction is rising. Organizations that cannot govern AI agents safely are accumulating compounding risks across three dimensions: operational, legal, and technical.

Pressure Points Driving Urgency

- Manual operations create latency, inconsistency, and headcount pressure
 - AI pilots without guardrails create legal and security risk
 - Tool sprawl increases operational complexity and governance burden
 - Enterprises need a standard operating layer for trusted agent execution
-

Product Overview

 Aligns with: Slide 4 – Product Overview

AI Agent Orchestration is a security-first SaaS control plane that allows enterprises to deploy, manage, and govern AI agents at scale — without rebuilding their existing workflows or accepting unacceptable security tradeoffs.

Core Capabilities

What the Platform Delivers

- Centralized orchestration, policy checks, scheduling, and telemetry
 - Multi-tenant design with per-tenant runtime isolation
 - Agent lifecycle intelligence: reuse idle capacity or spawn on demand
 - Web and Jira-based task intake with structured execution templates
-

Architecture & Trust Boundaries

 Aligns with: Slide 5 – Architecture

The platform operates on a shared control plane with fully isolated tenant execution environments. Every tenant workload runs in a dedicated AWS micro-account, providing hardware-level separation between customers and preventing any cross-tenant data exposure.

Architecture Highlights

Infrastructure Design

- Tenant workloads run in AWS micro-accounts (sub-accounts off payer)
- EKS runtime per tenant boundary for workload isolation
- Identity context retrieved and validated before execution
- Control decisions and logs centralized for governance visibility

Identity & Access Model

 *Aligns with: Slide 6 – Identity and Access Model*

Identity is enforced pre-execution, not retrofitted after an incident. This is a foundational design principle — every agent action is scoped to a verified identity before it runs, ensuring full traceability and minimizing the attack surface.

Identity Enforcement

- Scoped identities and least-privilege permissions per run
- Short-lived credentials to minimize standing access risk
- Policy gates to block unauthorized task actions
- Clear who/what/when traceability for each agent activity

Security & Compliance Posture

 *Aligns with: Slide 7 – Security and Compliance*

The platform's control stack is designed specifically for Fortune 500 operational requirements. Rather than treating compliance as a checkbox, the architecture embeds governance and auditability into every layer of execution.

Compliance-Ready Controls

- PII isolation through tenant-segmented data boundaries
- Private network controls and transport protections

- Continuous vulnerability scanning across agent runtime surfaces
 - Integrated patch lifecycle for agent workloads on EKS
 - Evidence-ready logging for internal governance and external audits
-

Workflow & User Experience

 *Aligns with: Slide 8 – Workflow and User Experience*

Enterprise teams submit work through systems they already use. There is no new behavior required — the platform meets users in Jira, web forms, and existing operational toolchains, then handles orchestration transparently behind the scenes.

Execution Flow

Step-by-Step Task Lifecycle

- Step 1: Task entered via web form or Jira issue
 - Step 2: Mandatory template defines scope, constraints, and expected outputs
 - Step 3: Orchestrator validates identity, policy, and runtime readiness
 - Step 4: Task routed to idle or newly provisioned agent
 - Step 5: Result returned with execution logs and status artifacts
-

On-Demand Scheduling

 *Aligns with: Slide 9 – On-Demand Scheduling Product*

A single platform handles both routine automation and burst operational demand. The scheduling engine manages resource allocation dynamically — routing work by policy, priority, and SLA — while providing full cost transparency for enterprise chargeback models.

Scheduling Capabilities

- Scheduled jobs for recurring high-volume workflows
 - Event-triggered workloads for urgent incidents and demand spikes
 - Resource orchestration by policy, priority, and SLA
 - Usage metering enables cost transparency and chargeback
-

Customer Profile & Initial Wedge

 Aligns with: Slide 10 – Customer Profile

The platform targets organizations where governance constraints are strongest and the cost of getting it wrong is highest — regulated industries, large enterprises, and security-sensitive environments where current AI tooling simply cannot pass procurement.

Ideal Customer Profile

Who We Sell To

- ICP: Fortune 500 and regulated enterprises
- Economic buyers: platform and security leadership
- Technical champions: DevOps, SRE, security engineering, enterprise architecture
- Wedge use cases: incident enrichment, ticket triage, patch verification

Business Model

 Aligns with: Slide 11 – Business Model

Revenue is designed to be predictable at the platform level and expand naturally with usage. Enterprise customers subscribe for access to the control plane, then scale billing organically as they run more workflows across more teams.

Revenue Structure

- Annual subscription for control plane and baseline enterprise controls
- Usage-based billing for execution volume and on-demand workloads
- Optional enterprise modules for advanced governance and networking
- Expansion dynamics: more workflows, teams, and workload intensity

Go-To-Market Strategy

 Aligns with: Slide 12 – Go-To-Market Strategy

The GTM motion is compliance-led and integration-first. Rather than competing on agent features, the platform lands through security and risk conversations — where procurement cycles are clearer, budgets are allocated, and urgency is felt most acutely.

GTM Phases

- Land through design partners in security-sensitive organizations
- Prove value in 60-90 day pilots against risk and productivity KPIs
- Expand from one workflow to multi-team standardization
- Channel leverage via cloud/security partners and integration ecosystem

Competitive Positioning

 Aligns with: Slide 13 – Competitive Positioning

The platform does not compete on raw agent capability — it competes on operational trust. This is a deliberate positioning choice that targets a less crowded, higher-value segment of the market: enterprises that cannot move forward without governance infrastructure.

Structural Differentiators

- Account-level tenant isolation architecture (not shared runtimes)
- Identity and policy validation before execution (not post-incident)
- Integrated security scan and patch lifecycle in runtime
- Structural moat: trust architecture + orchestration intelligence + workflow embedding

Roadmap & Milestones

 Aligns with: Slide 14 – Roadmap and Milestones

Execution is organized across three phases, each building on the prior to deliver compounding value: hardening the product, expanding integrations, and scaling deployment velocity.

0–6 Months: Pilot & Hardening

- Pilot delivery with design partner customers
- Production hardening and control depth
- Core governance and audit infrastructure

6–12 Months: Integration & Automation

- Integration expansion across enterprise toolchains
- Governance automation features
- Initial enterprise conversion from pilots

12–18 Months: Scale & Deployment

- Onboarding velocity improvements
- Broader enterprise deployment playbooks
- Success metrics: paid pilots, annual conversions, growing workload volume

Utility Token for Autonomous Workloads

 *Aligns with: Slide 14b – Token Model Consideration*

As AI agents move from assisted automation toward fully autonomous operation, a native utility token presents a compelling mechanism for metering, incentivizing, and governing workload execution at scale. Rather than a speculative financial instrument, the token would function as the economic unit of the agent execution layer — consumed when work is done, earned when infrastructure is contributed.

How It Works

Each autonomous workload consumes tokens proportional to its execution complexity, resource usage, and SLA tier. Enterprises purchase token allocations as part of their subscription or on-demand — giving finance teams predictable cost controls and giving platform operators a transparent utilization signal.

Token Utility Functions

- Execution credits: tokens are spent to initiate and sustain autonomous agent runs
- Priority routing: higher token stakes unlock faster scheduling and dedicated capacity
- Governance staking: operators stake tokens to assert policy compliance, creating accountability
- Ecosystem access: tokens gate access to specialized agent models and third-party integrations
- Chargeback transparency: token ledger provides auditable per-team cost attribution

Why This Fits the Platform

The orchestration layer already tracks every execution event — task origin, resource consumption, identity, duration, and outcome. This telemetry is the natural foundation for a token economy. The platform does not need to retrofit token logic; it already has the instrumentation required to make it work.

Strategic Considerations

- Token model is additive — the core SaaS business stands independently
- Regulatory clarity is required before any public token issuance
- Initial deployment could be permissioned (enterprise-only) before any broader rollout
- Token design should prioritize utility and compliance over yield or speculation
- Aligns platform incentives with customer outcomes: more successful runs, more tokens consumed

This remains an active area of design exploration. The decision to formalize a token model will be informed by enterprise customer feedback, regulatory developments, and the maturity of the execution telemetry layer.

Fundraise & Use of Funds

 Aligns with: Slide 15 – Fundraise and Use of Funds

Seed capital will be deployed to establish category leadership in secure agent execution — accelerating enterprise readiness, building repeatable GTM, and positioning the company as the infrastructure standard for trusted AI operations.

Use of Funds Allocation

- 60% — Product and security engineering
- 25% — Go-to-market and sales
- 15% — Operations and compliance

Primary goals: accelerate enterprise readiness, establish GTM repeatability, and unlock expansion growth. Investor outcome: exposure to durable enterprise AI infrastructure — a category that will only grow as autonomous execution becomes standard.

Diligence Appendix

The following materials are available upon request for second meetings, technical diligence, and partner deep dives:

- Security control matrix mapped to buyer questionnaire categories
- Tenant trust boundary diagrams with data flow breakdown
- Pilot blueprint with entry criteria, success metrics, and conversion triggers
- Hiring plan aligned to roadmap and go-to-market milestones

For access to diligence materials or to schedule a technical deep dive, please contact the founding team directly.